

---

```

from __future__ import print_function
import numpy as np
import argparse
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms
from sklearn.metrics import *
from matplotlib import pyplot as plt
%matplotlib inline

class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Conv2d(1, 20, 5, 1)
        self.conv2 = nn.Conv2d(20, 50, 5, 1)
        self.fc1 = nn.Linear(4*4*50, 500)
        self.fc2 = nn.Linear(500, 10)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = F.max_pool2d(x, 2, 2)
        x = F.relu(self.conv2(x))
        x = F.max_pool2d(x, 2, 2)
        x = x.view(-1, 4*4*50)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return F.log_softmax(x, dim=1)

def train(model, device, train_loader, optimizer, epoch):
    losses = []
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        output = model(data)
        loss = F.nll_loss(output, target)
        loss.backward()
        optimizer.step()
        losses.append(loss.item())
        if batch_idx > 0 and batch_idx % 100 == 0:
            print('Train Epoch: {} [{}/{}]\t({:.0f}%) \tLoss: {:.6f}'.format(
                epoch, batch_idx * len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.item()))
    return losses

def test(model, device, test_loader):
    model.eval()
    test_loss = 0
    correct = 0
    with torch.no_grad():
        for data, target in test_loader:
            data, target = data.to(device), target.to(device)

```

```

        data, target = data_loader(device), target_loader(device),
        output = model(data)
        test_loss += F.nll_loss(output, target, reduction='sum').item() # sum up batch
        pred = output.argmax(dim=1, keepdim=True) # get the index of the max log-proba
        correct += pred.eq(target.view_as(pred)).sum().item()
    test_loss /= len(test_loader.dataset)
    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%) \n'.format(
        test_loss, correct, len(test_loader.dataset),
        100. * correct / len(test_loader.dataset)))
    return (float(correct) / len(test_loader.dataset))

train_loader = torch.utils.data.DataLoader(
    datasets.MNIST(
        '../data',
        train=True,
        download=True,
        transform=transforms.Compose([
            transforms.ToTensor(),
            transforms.Normalize((0.1307,), (0.3081,))
        ])
    ),
    batch_size=64,
    shuffle=True)
test_loader = torch.utils.data.DataLoader(
    datasets.MNIST(
        '../data',
        train=False,
        transform=transforms.Compose([
            transforms.ToTensor(),
            transforms.Normalize((0.1307,), (0.3081,))
        ])
    ),
    batch_size=1000,
    shuffle=True)

model = CNN()
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.5)
device = torch.device("cpu") # or 'gpu'
losses = []
accuracies = []
for epoch in range(0, 10):
    losses.extend(train(model, device, train_loader, optimizer, epoch))
    accuracies.append(test(model, device, train_loader))

```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>  
Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to  
100% 9912422/9912422 [00:00<00:00,  
7680848.56it/s]

Extracting ../data/MNIST/raw/train-images-idx3-ubyte.gz to ../data/MNIST/r

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz>  
Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to  
100% 28881/28881 [00:00<00:00,  
554188.79it/s]

Extracting ../data/MNIST/raw/train-labels-idx1-ubyte.gz to ../data/MNIST/r

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz>  
Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to  
100% 1648877/1648877 [00:00<00:00,  
7460944.47it/s]

Extracting ../data/MNIST/raw/t10k-images-idx3-ubyte.gz to ../data/MNIST/r

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz>  
Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to  
100% 4542/4542 [00:00<00:00,  
126601.78it/s]

Extracting ../data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ../data/MNIST/r

Train Epoch: 0	[6400/60000	(11%)]	Loss: 0.521960
Train Epoch: 0	[12800/60000	(21%)]	Loss: 0.325841
Train Epoch: 0	[19200/60000	(32%)]	Loss: 0.203964
Train Epoch: 0	[25600/60000	(43%)]	Loss: 0.136594
Train Epoch: 0	[32000/60000	(53%)]	Loss: 0.236310
Train Epoch: 0	[38400/60000	(64%)]	Loss: 0.274870
Train Epoch: 0	[44800/60000	(75%)]	Loss: 0.215008
Train Epoch: 0	[51200/60000	(85%)]	Loss: 0.083802
Train Epoch: 0	[57600/60000	(96%)]	Loss: 0.197037

Test set: Average loss: 0.1080, Accuracy: 58013/60000 (97%)

Train Epoch: 1	[6400/60000	(11%)]	Loss: 0.039530
Train Epoch: 1	[12800/60000	(21%)]	Loss: 0.228269
Train Epoch: 1	[19200/60000	(32%)]	Loss: 0.140882
Train Epoch: 1	[25600/60000	(43%)]	Loss: 0.056366
Train Epoch: 1	[32000/60000	(53%)]	Loss: 0.042062
Train Epoch: 1	[38400/60000	(64%)]	Loss: 0.056974
Train Epoch: 1	[44800/60000	(75%)]	Loss: 0.021462
Train Epoch: 1	[51200/60000	(85%)]	Loss: 0.102775
Train Epoch: 1	[57600/60000	(96%)]	Loss: 0.046871

Test set: Average loss: 0.0622, Accuracy: 58870/60000 (98%)

Train Epoch: 2	[6400/60000	(11%)]	Loss: 0.147645
Train Epoch: 2	[12800/60000	(21%)]	Loss: 0.093769
Train Epoch: 2	[19200/60000	(32%)]	Loss: 0.083718
Train Epoch: 2	[25600/60000	(43%)]	Loss: 0.035652
Train Epoch: 2	[32000/60000	(53%)]	Loss: 0.013875
Train Epoch: 2	[38400/60000	(64%)]	Loss: 0.032025

Train Epoch: 2 [44800/60000 (75%)] Loss: 0.031964  
Train Epoch: 2 [51200/60000 (85%)] Loss: 0.047817  
Train Epoch: 2 [57600/60000 (96%)] Loss: 0.036271

Test set: Average loss: 0.0511, Accuracy: 59061/60000 (98%)

Train Epoch: 3 [6400/60000 (11%)] Loss: 0.017703  
Train Epoch: 3 [12800/60000 (21%)] Loss: 0.014950  
Train Epoch: 3 [19200/60000 (32%)] Loss: 0.353113  
Train Epoch: 3 [25600/60000 (43%)] Loss: 0.051852  
Train Epoch: 3 [32000/60000 (53%)] Loss: 0.013181  
Train Epoch: 3 [38400/60000 (64%)] Loss: 0.012055  
Train Epoch: 3 [44800/60000 (75%)] Loss: 0.005071  
Train Epoch: 3 [51200/60000 (85%)] Loss: 0.011054  
Train Epoch: 3 [57600/60000 (96%)] Loss: 0.016299

Test set: Average loss: 0.0405, Accuracy: 59258/60000 (99%)

Train Epoch: 4 [6400/60000 (11%)] Loss: 0.146109  
Train Epoch: 4 [12800/60000 (21%)] Loss: 0.011323  
Train Epoch: 4 [19200/60000 (32%)] Loss: 0.002555  
Train Epoch: 4 [25600/60000 (43%)] Loss: 0.185082  
Train Epoch: 4 [32000/60000 (53%)] Loss: 0.046304  
Train Epoch: 4 [38400/60000 (64%)] Loss: 0.126782  
Train Epoch: 4 [44800/60000 (75%)] Loss: 0.018411  
Train Epoch: 4 [51200/60000 (85%)] Loss: 0.097465  
Train Epoch: 4 [57600/60000 (96%)] Loss: 0.050512

Test set: Average loss: 0.0335, Accuracy: 59388/60000 (99%)

Train Epoch: 5 [6400/60000 (11%)] Loss: 0.027327  
Train Epoch: 5 [12800/60000 (21%)] Loss: 0.026542  
Train Epoch: 5 [19200/60000 (32%)] Loss: 0.018746  
Train Epoch: 5 [25600/60000 (43%)] Loss: 0.012916  
Train Epoch: 5 [32000/60000 (53%)] Loss: 0.062889  
Train Epoch: 5 [38400/60000 (64%)] Loss: 0.004592  
Train Epoch: 5 [44800/60000 (75%)] Loss: 0.020959  
Train Epoch: 5 [51200/60000 (85%)] Loss: 0.100499  
Train Epoch: 5 [57600/60000 (96%)] Loss: 0.006967

Test set: Average loss: 0.0264, Accuracy: 59534/60000 (99%)

Train Epoch: 6 [6400/60000 (11%)] Loss: 0.024826  
Train Epoch: 6 [12800/60000 (21%)] Loss: 0.032037  
Train Epoch: 6 [19200/60000 (32%)] Loss: 0.003827  
Train Epoch: 6 [25600/60000 (43%)] Loss: 0.001596  
Train Epoch: 6 [32000/60000 (53%)] Loss: 0.023506  
Train Epoch: 6 [38400/60000 (64%)] Loss: 0.022206  
Train Epoch: 6 [44800/60000 (75%)] Loss: 0.054582  
Train Epoch: 6 [51200/60000 (85%)] Loss: 0.008049  
Train Epoch: 6 [57600/60000 (96%)] Loss: 0.038125

Test set: Average loss: 0.0281, Accuracy: 59492/60000 (99%)

Train Epoch: 7 [6400/60000 (11%)] Loss: 0.039052  
Train Epoch: 7 [12800/60000 (21%)] Loss: 0.007173  
Train Epoch: 7 [19200/60000 (32%)] Loss: 0.007180  
Train Epoch: 7 [25600/60000 (43%)] Loss: 0.026620  
Train Epoch: 7 [32000/60000 (53%)] Loss: 0.006726  
Train Epoch: 7 [38400/60000 (64%)] Loss: 0.003266  
Train Epoch: 7 [44800/60000 (75%)] Loss: 0.001000  
Train Epoch: 7 [51200/60000 (85%)] Loss: 0.001000  
Train Epoch: 7 [57600/60000 (96%)] Loss: 0.001000

```
Train Epoch: 7 [44800/60000      (75%)] Loss: 0.011263
Train Epoch: 7 [51200/60000      (85%)] Loss: 0.049789
Train Epoch: 7 [57600/60000      (96%)] Loss: 0.077939
```


Test set: Average loss: 0.0229, Accuracy: 59589/60000 (99%)

```
Train Epoch: 8 [6400/60000       (11%)] Loss: 0.006500
Train Epoch: 8 [12800/60000      (21%)] Loss: 0.011150
Train Epoch: 8 [19200/60000      (32%)] Loss: 0.008179
Train Epoch: 8 [25600/60000      (43%)] Loss: 0.057078
Train Epoch: 8 [32000/60000      (53%)] Loss: 0.016232
Train Epoch: 8 [38400/60000      (64%)] Loss: 0.013455
Train Epoch: 8 [44800/60000      (75%)] Loss: 0.008888
Train Epoch: 8 [51200/60000      (85%)] Loss: 0.034871
Train Epoch: 8 [57600/60000      (96%)] Loss: 0.004731
```

Test set: Average loss: 0.0271, Accuracy: 59465/60000 (99%)


```
def mean(li): return sum(li)/len(li)
plt.figure(figsize=(14, 4))
plt.xlabel('training batch')
plt.ylabel('loss')
plt.plot([mean(losses[i:i+10]) for i in range(len(losses))])
```

[<matplotlib.lines.Line2D at 0x7fb6e9234b10>]



```
plt.figure(figsize=(14, 4))
plt.xticks(range(len(accuracies)))
plt.xlabel('training epoch')
plt.ylabel('accuracy')
plt.plot(accuracies, marker='o')
```

[<matplotlib.lines.Line2D at 0x7fb6e92c1990>]



```
def test_label_predictions(model, device, test_loader):
    model.eval()
    actuals = []
    predictions = []
    with torch.no_grad():
        for data, target in test_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            prediction = output.argmax(dim=1, keepdim=True)
            actuals.extend(target.view_as(prediction))
            predictions.extend(prediction)
    return [i.item() for i in actuals], [i.item() for i in predictions]

actuals, predictions = test_label_predictions(model, device, test_loader)
print('Confusion matrix:')
print(confusion_matrix(actuals, predictions))
print('F1 score: %f' % f1_score(actuals, predictions, average='micro'))
print('Accuracy score: %f' % accuracy_score(actuals, predictions))
```

Confusion matrix:

```
[[ 976    0    0    0    0    0    1    1    2    0]
 [    0 1131    2    0    0    1    1    0    0    0]
 [    1    0 1030    0    0    0    0    1    0    0]
 [    0    0    2 1004    0    3    0    0    1    0]
 [    1    0    2    0 961    0    2    1    1   14]
 [    2    0    0    6    0 877    1    1    2    3]
 [    2    1    0    1    1    6 947    0    0    0]
 [    0    3    8    0    0    0    0 1010    2    5]
 [    1    0    2    3    0    0    1    1 963    3]
 [    1    2    0    1    1    2    0    0    2 1000]]
```

F1 score: 0.989900

Accuracy score: 0.989900

```
def test_class_probabilities(model, device, test_loader, which_class):
    model.eval()
    actuals = []
    probabilities = []
    with torch.no_grad():
        for data, target in test_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            prediction = output.argmax(dim=1, keepdim=True)
            actuals.extend(target.view_as(prediction) == which_class)
            probabilities.extend(np.exp(output[:, which_class]))
    return [i.item() for i in actuals], [i.item() for i in probabilities]
```


which\_class = 9

```

actuals, class_probabilities = test_class_probabilities(model, device, test_loader, which_

fpr, tpr, _ = roc_curve(actuals, class_probabilities)
roc_auc = auc(fpr, tpr)
plt.figure()
lw = 2
plt.plot(fpr, tpr, color='darkorange',
         lw=lw, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC for digit=%d class' % which_class)
plt.legend(loc="lower right")
plt.show()

```



```

print('Trainable parameters:')
for name, param in model.named_parameters():
    if param.requires_grad:
        print(name, '\t', param.numel())

```

```

Trainable parameters:
conv1.weight      500
conv1.bias        20
conv2.weight      25000
conv2.bias        50
fc1.weight        400000
fc1.bias          500
fc2.weight        5000
fc2.bias          10

```

